

# System Design

## Agentic Orchestration Builder: System Design Document

Build an event-driven, stateful orchestration platform that unifies deterministic DAGs and adaptive multi-agent workflows under a single control plane, with durable execution, human-in-the-loop checkpoints, and full replay/observability. Ground the runtime in durable workflow primitives (e.g., Temporal/Zeebe), overlay agentic coordination patterns for non-determinism, and expose a drag-and-drop builder plus AI copilot for design, governance, and execution.

### 1) Problem Statement and Vision

Modern automations must handle both:

- Deterministic flows: predictable, repeatable DAGs with clear success/error paths.
- Non-deterministic flows: agentic, context-driven decisions, dynamic branching, reflection/critique loops, and human approvals at critical steps.

The platform must:

- Treat events as the control plane to drive orchestration, communication, and recovery.
- Provide durable state, replay, timeouts, retries, and compensations.
- Natively support agentic patterns (coordinator, hierarchical, parallel/concurrent, iterative refinement, human-in-the-loop) with guardrails and approvals.
- Offer a visual builder and an AI co-pilot that elicits requirements and co-designs robust, secure, cost-aware workflows.

### 2) Functional Requirements

## 2.1 Workflow Authoring and Execution

- Visual drag-and-drop builder for triggers, actions, conditions, parallel branches, loops, compensations, and agent steps (deterministic nodes and agentic nodes).
- Import/export BPMN/DMN (optional) for enterprise process teams.
- Event-driven triggers (webhooks, message topics, schedules, file/object events, DB CDC).
- Deterministic execution with durable timers, retries, idempotency, compensations, and saga patterns.
- Non-deterministic/agentic execution with coordinator/manager agents, reflection loops, and human checkpoints.

## 2.2 Agentic Capabilities

- Multi-agent coordinator to decompose goals and dispatch tasks to specialized agents; supports hierarchical decomposition and concurrent fan-out/fan-in.
- Tool use library (HTTP/API, DB, RPA connectors, scraping, email, chat, files), with scoped credentials per agent (RBAC).
- Context engineering and memory: vector store for embeddings, episodic/execution memory, secure context windows, and per-tenant/flow isolation.
- Safety and governance: policy checks, rate/cost limits, approval gates, risk scoring, and rollback/containment strategies.

## 2.3 Human-in-the-Loop

- Inline approval, multi-party review, escalation and SLAs, with durable pause/resume and replay.
- "Critical step confirmation" prompts injected by AI copilot at risk-sensitive nodes, configurable by policy.

## 2.4 Event and State Management

- Event ingestion (webhooks, connectors), enrichment, routing to workflows/agents, dead-letter queues, and ordered/partitioned processing.

- Durable state store for workflow definitions, executions, decisions, timers, and checkpoints; full audit, replay, and time travel.

## 2.5 AI Copilot (Automation Mode)

- Requirement elicitation: interrogates user for intent, scope, security, data sensitivity, SLAs, costs, margins, ROI.
- Co-architects workflows (deterministic vs agentic), proposes nodes and guardrails, inserts approvals at critical points, and estimates costs/latency.
- Auto-generates test data, simulators, and golden-path/equivocation tests for reliability.

## 2.6 Integrations and Connectors

- Out-of-the-box connectors: HTTP, Webhooks, DBs, S3/GCS, email/SMS/WhatsApp, calendars, Slack/Teams, CRMs/ERPs, RPA bridges.
- LLM backends: OpenAI/Anthropic/local via adapters for cost/latency control; routing to faster/smaller models for simpler tools.

## 2.7 Observability and Analytics

- End-to-end tracing with execution graphs, per-node timings, token usage, cost per run, failure clusters, drift detection.
- Policy/reporting dashboards: SLA attainment, approval latencies, cost centres, model/tool health.

## 2.8 Multi-tenancy and Enterprise Controls

- Hard isolation by tenant with quotas, rate limits, data residency and region pinning, and org-level RBAC/policies.

## 2.9 Marketplace and Templates

- Reusable node packs, agent roles, starter blueprints for industry use-cases. Certification and signed bundles.

# 3) Non-Functional Requirements

Category	Target/Requirement
API read latency	P99 < 200 ms

API write latency	P99 < 500 ms
Webhook ingress	< 100 ms processing
Deterministic step latency	< 100 ms median (non-external)
Agentic decision step	0.5–3 s depending on model/tool
Availability	99.9% baseline; HA across AZs; graceful degradation on LLM/provider failures
Durability	11 nines for logs/artifacts via replicated object storage; DB multi-AZ; WAL/backups
Security/Compliance	TLS 1.3, AES-256 at rest, secrets manager, RBAC/ABAC, audit trails, PII detection/masking; SOC2/GDPR-ready
Data residency	Tenant-pinned stores and compute with regionalization
Cost efficiency	Autoscaling, model routing, caching, batch where tolerable; robust cost telemetry

## 4) Back-of-the-Envelope Estimates

Assumptions (initial public launch):

- 50,000 MAU; 10 workflows/user; 2M executions/day; 10M agent invocations/day; 5M AI API calls/day; peak/avg QPS  $\approx$  5x.
- Storage growth:  $\sim$ 500 GB/month for execution logs and artifacts (object storage, compressed).
- Peak QPS targets: 5,000 for API+event ingress; DB IOPS target  $\sim$ 15,000 with 85% cache hit.

Generated capacity/cost CSV assets

Interpretation:

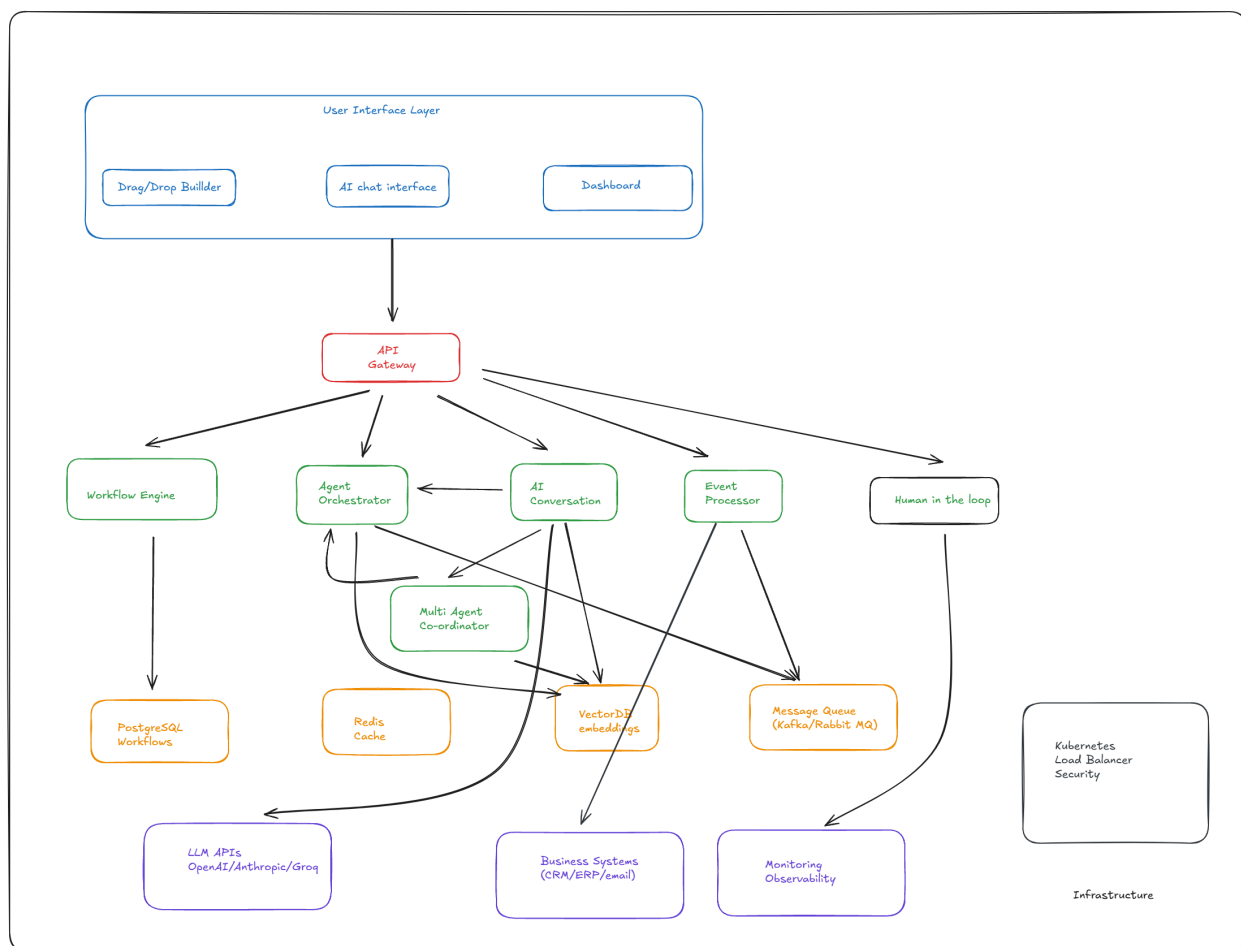
- Kafka partitions sized for 5–10K msg/s with 3x replication; consumer groups per service.
- PostgreSQL with read replicas, connection pooling; consider sharding by tenant\_id at  $>1K$  tenants per shard benchmark.
- Redis for session/coordination; TTL strategies to keep cache hit  $\geq 85\%$  and reduce DB IOPS.

- Vector DB sized for prompts/memories; per-tenant namespaces; tiered retention.

Rough monthly costs (cloud-managed, blended; varies by region/provider):

- Compute (Kubernetes): \$8K; DB: \$3.5K; Cache: \$0.8K; Storage: \$1.2K; Kafka: \$0.6K; LLM: \$15K; Vector DB: \$1.5K; Monitoring: \$0.5K; LB/DNS: \$0.4K → ~\$31.5K/month at this scale.

## 5) High-Level Architecture



- UI: Drag-and-drop builder; AI chat copilot; Execution dashboards.
- API Gateway: AuthN/Z, rate limiting, request validation, tenant routing.
- Core Services:

- Workflow Engine (deterministic): durable execution, timers, retries, compensations (Temporal/Zeebe).
- Event Processing: webhook ingress, enrichment, routing, DLQs, replays; Kafka topics and schemas.
- Multi-Agent Coordinator: implements coordinator/hierarchical/concurrent patterns; task ledger; cost/latency budgeter; tool adapters.
- AI Conversation Service: requirement elicitation, planning, critique/reflection, grounding; prompt libraries; secure toolformer calls.
- Human-in-the-Loop: approval tasks, escalations, SLAs; durable pause/resume; policy injection.
- Policy & Governance: rules, risk scoring, PEM/OPA evaluation, secrets and vault integration.
- Connector Hub: outbound/inbound connectors with backpressure, rate controls, token buckets, and provider-specific retry semantics.
- Data Plane:
  - PostgreSQL: tenants, users, workflow\_defs, executions, step\_executions, approvals, policies, credentials metadata.
  - Redis: sessions, hot keys, distributed locks, short-lived coordination.
  - Vector DB: embedding stores for contexts, memories, tool schemas; per-tenant namespaces.
  - Kafka: event buses per domain (triggers, approvals, agent-tasks, audits, DLQ).
  - Object Storage: artifacts, transcripts, audit bundles, replays.
- Infrastructure:
  - Kubernetes: autoscaling, HPA/VPA, node pools (general, CPU-bound, GPU-bound for local inference).
  - Observability: Prometheus/Grafana, tracing, ELK; SLOs and alerting playbooks.
  - Security: mTLS inter-service, secret management, KMS; egress controls.

Summary:

- Event-driven backbone with Kafka/EventBridge-equivalent for decoupling.
- Deterministic durability via workflow engine (Temporal/Zeebe/iWF).
- Agentic orchestration patterns (coordinator, hierarchical, concurrent, iterative refinement, human-in-the-loop) per cloud/industry references.
- Drag-and-drop builder usability principles.

## **6) Database Design (core tables)**

Tenancy and Identity

- tenants(id, region, plan, data\_residency)
- users(id, tenant\_id, email, role, mfa, status)
- service\_accounts(id, tenant\_id, scopes)
- api\_keys(id, tenant\_id, hash, rate\_limit, expires\_at)

Workflow Model

- workflow\_definitions(id, tenant\_id, name, version, spec\_json, type: deterministic|agentic|hybrid, active, created\_at, created\_by)
- workflow\_executions(id, workflow\_id, tenant\_id, status, current\_step, started\_at, completed\_at, context\_json, cost\_cents)
- step\_executions(id, execution\_id, step\_index, step\_type, status, input\_json, output\_json, error, started\_at, completed\_at)
- schedules(id, tenant\_id, cron, workflow\_id, enabled)
- webhooks(id, tenant\_id, source, secret, filters\_json, enabled)

Agentic Model

- agents(id, tenant\_id, role, tools\_json, policy\_json, status, last\_active\_at)
- agent\_tasks(id, tenant\_id, agent\_id, kind, priority, payload\_json, status, scheduled\_at, started\_at, completed\_at, result\_json)
- task\_ledger(id, execution\_id, parent\_id, title, state, assigned\_agent\_id, cost\_budget\_cents, token\_budget, reason\_json)

## Approvals and Policy

- approval\_tasks(id, tenant\_id, execution\_id, step\_id, approvers\_json, sla\_seconds, status, decided\_by, decided\_at, decision, comment)
- policies(id, tenant\_id, name, type, rules\_json, enabled)
- secrets\_metadata(id, tenant\_id, name, scope, rotated\_at)

## Events and Audit

- events(id, tenant\_id, type, source, payload\_json, correlation\_id, status, created\_at, processed\_at)
- audits(id, tenant\_id, actor\_type, actor\_id, action, resource\_type, resource\_id, diff\_json, ts)

## Analytics/Quotas

- usage\_counters(id, tenant\_id, metric, window\_start, window\_end, count, cost\_cents)
- sla\_breaches(id, tenant\_id, execution\_id, kind, details\_json, ts)

## Vector Store

- Stored in external vector DB with metadata: tenant\_id, workflow\_id, execution\_id, type (memory, tool\_doc, KB), embedding, pointer\_uri.

## Indexes and partitioning:

- Composite indexes for tenant\_id + status/stage columns; time-partition large append-only tables (events, audits, step\_executions) for pruning.
- Consider sharding by tenant\_id for very large tenancy cohorts.

# 7) API Design (some examples)

## Auth

- POST /v1/auth/token
- GET /v1/me

## Workflows

- GET /v1/workflows?tenant=...&active=true



- POST /v1/workflows {name, type, spec}
- POST /v1/workflows/{id}/execute {input, correlation\_id}
- GET /v1/executions/{id}
- POST /v1/webhooks/ingest (HMAC-verified)

#### Agentic

- POST /v1/agents {role, tools, policy}
- POST /v1/agents/{id}/tasks {payload}
- GET /v1/executions/{id}/ledger
- POST /v1/executions/{id}/reflect (trigger critique/replan)

#### Human-in-the-Loop

- GET /v1/approvals?status=pending
- POST /v1/approvals/{id}/decision {approve|reject, comment}

#### Governance

- POST /v1/policies {rules}
- GET /v1/costs?window=...

#### Observability

- GET /v1/executions/{id}/graph
- GET /v1/metrics/executions?group\_by=workflow
- GET /v1/traces/{trace\_id}

Webhook security: HMAC signature, replay windows, IP allowlists.

Pagination: cursor-based; idempotency keys for create/execute.

Rate limits: per-tenant tier; burst allowances; per-key quotas.

## 8) Orchestration Patterns Supported

- Deterministic: sequential, parallel fan-out/fan-in, loop, compensation/saga, timeouts; best for predictable tasks.

- Coordinator pattern: AI manager routes tasks to specialized agents; adaptive branching; tool-based actions.
- Hierarchical task decomposition: multi-level planning for ambiguous problems, with higher cost/latency trade-offs.
- Concurrent orchestration: multiple agents in parallel with aggregation and selection.
- Iterative refinement and review/critique loops: quality gates and self-correction.
- Human-in-the-loop: approval and escalation pattern; policy-driven checklists.

## 9) Execution Semantics and Safety

Deterministic steps:

- Pure functions with retries and backoff; idempotent side effects; compensations for non-idempotent actions; exactly-once intent with outbox pattern.

Agentic steps:

- Budgeted token/cost ceilings; timeouts; tool call whitelists; sandboxed outputs; pre/post validation hooks; red-teaming prompts; auto-insert approval on high-risk scopes.
- Fall-back strategies: smaller/cheaper models, cached responses, human escalation.

Event handling:

- Ordering where needed via partition keys; DLQs with reprocess UI; poison-pill detection; circuit breakers.

## 10) Observability and SLOs

Metrics

- Request rate/latency/error per API; execution duration per workflow; step distribution; queue lag; cache hit; DB latency; token/cost per execution; model/provider latency.

### Tracing

- Distributed tracing with correlation IDs across workflow, agent, tool calls.

### Logging

- Structured logs with redaction for PII; long-term archival tiers with lifecycle policies.

### SLOs (starter)

- API availability 99.9%; execution success rate  $\geq 99\%$  for deterministic; mean approval latency SLA per policy; webhook processing P95 < 100 ms.

### Alerts

- Error rate spikes; queue lag; approval SLA breaches; cost budget overruns; model timeouts and provider saturation.

## 11) Security, Compliance, and Governance

- RBAC/ABAC for users, service accounts, and agents; resource-scoped permissions and allowlisted tools.
- Data protection: encryption in transit and at rest, field-level encryption for sensitive fields, PII detection and masking.
- Tenant isolation: network and data segregation; per-tenant encryption contexts; region pinning; key rotation.
- Governance and audit: immutable audit trails for user/agent actions, prompt and tool logs, replay artifacts; policy packs per vertical.

## 12) Deployment and Scaling

- Kubernetes with autoscaling; distinct pools for CPU vs GPU; canary/blue-green with metrics gates.
- Horizontal scaling by microservice; Kafka partitions sized to traffic domains; Redis clustering; DB read replicas; shard by tenant at scale.
- Multi-region: active-active stateless, active-passive for DBs; cross-region replication with bounded staleness; tenant-level residency controls.

## **13) Migration and Reliability Practices**

- Zero-downtime schema migrations (expand/migrate/contract); dual writes and backfills with validation.
- Chaos testing in staging for workflow timeouts, DLQs, and agentic fallbacks.
- Replay tooling from archived events with deterministic vs agentic semantic flags.

## **14) Cost Optimization**

- Model routing/caching; response reuse for deterministic prompts; toolformer-style planner to minimize LLM calls; summarize/condense contexts.
- Selective tracing/logging levels; lifecycle policies for storage; spot/interruptible nodes for batch.
- Connector rate shaping and batching to cut provider costs.
- Per-tenant cost dashboards and budgets with automatic throttling/escalation.

## **15) Product UX Highlights**

- Builder: minimal-text, visual-first canvas with node packs (deterministic, agent, human, event, compensations), inline cost/latency estimates per path, and policy hints.
- Copilot: requirements interview, drafts initial workflow, highlights critical steps needing confirmation, simulates runs with synthetic event streams (no PII), recommends guardrails and tests.
- Execution view: live swimlanes (agents, human, systems), replay slider, and "what-if" branching.
- Templates: industry blueprints for sales ops, support triage, finance approvals, supply chain exceptions.